Handout 2: Gradient descent

Lecturer & author: Georgios P. Karagiannis   georgios.karagiannis@durham.ac.uk

Aim. To introduce gradient descent, its motivation, description, practical tricks, analysis in the convex scenario, and implementation.

Reading list & references:

- Shalev-Shwartz, S., & Ben-David, S. (2014). Understanding machine learning: From theory to algorithms. Cambridge university press.
  - Ch. 14.1 Gradient Descent

Further reading

- Bishop, C. M. (2006). Pattern recognition and machine learning. New York: Springer.

1. MOTIVATIONS

Note 1. Consider a learning problem $(\mathcal{H}, \mathcal{Z}, \ell)$. Learning may involve the computation of the minimizer $h^* \in \mathcal{H}$, where $\mathcal{H}$ is a class of hypotheses, of the empirical risk function (ERF) $\hat{R}(h) = \frac{1}{n} \sum_{i=1}^{n} \ell(h, z_i)$ given a finite sample $\{z_i; i = 1, ..., n\}$ generated from the data generating model $g(\cdot)$ and using loss $\ell(\cdot)$; that is

$$(1.1) \qquad h^* = \arg \min_{\forall h \in \mathcal{H}} \left(\hat{R}(h) \right) = \arg \min_{\forall h \in \mathcal{H}} \left(\frac{1}{n} \sum_{i=1}^{n} \ell(h, z_i) \right)$$

If analytical minimization of (1.1) is impossible or impractical, numerical procedures can be applied; eg Gradient Descent (GD) algorithms. Such approaches introduce numerical errors in the solution.

2. DESCRIPTION

Notation 2. For the sake of notation simplicity and generalization, we will present Gradient Descent (GD) in the following minimization problem

$$(2.1) \qquad w^* = \arg \min_{\forall w \in \mathcal{H}} (f(w))$$

where here $f : \mathbb{R}^d \to \mathbb{R}$, and $w \in \mathcal{H} \subseteq \mathbb{R}^d$; $f(\cdot)$ is the function to be minimized, e.g., $f(\cdot)$ can be an empirical risk function $\hat{R}(\cdot)$.

Assumption 3. *Assume (for now) that $f : \mathbb{R}^d \to \mathbb{R}$ is a differentiable function.*

Definition 4. Given a per-specified learning rate $\eta_t > 0$, the Gradient Descent (GD) algorithm for the solution of the minimization problem (2.1) is given in Algorithm 1

Algorithm 1 Gradient descent algorithm with learning rate η_t

---

For $t = 1, 2, 3, \dots$ iterate:

    (1) compute

(2.2)
$$w^{(t+1)} = w^{(t)} - \eta_t \nabla f\left(w^{(t)}\right)$$

    (2) terminate if a termination criterion is satisfied, e.g.

$$\text{If } t \geq T_{\max} \text{ then STOP}$$

---

where

$$\nabla f(w) = \left(\frac{\partial}{\partial x_1} f(x), \dots, \frac{\partial}{\partial x_d} f(x)\right)^{\top} \Bigg|_{x=w}$$

is the gradient of f at w.

Remark 5. Interpreting GD, GD produces a chain $\left\{w^{(t)}\right\}$ that drifts towards the minimum w^*. It evolves directed towards the opposite direction than that of the gradient $\nabla f(\cdot)$ and at a rate controlled by the learning rate η_t.

Remark 6. For more intuitive explanation, consider the (1st order) Taylor polynomial for the approximation of $f(w)$ in a small area around u (i.e. $\|v - u\| = \text{small}$)

$$f(u) \approx P(u) = f(w) + \langle u - w, \nabla f(w)\rangle$$

Assuming convexity for f, it is

See

Handout 1

(2.3)
$$f(u) \geq \underbrace{f(w) + \langle u - w, \nabla f(w)\rangle}_{=P(u;w)}$$

meaning that P lower bounds f. Hence we could design an updating mechanism producing $w^{(t+1)}$ which is nearby $w^{(t)}$ (small steps) and which minimize the linear approximation $P(w)$ of $f(w)$ at $w^{(t)}$

(2.4)
$$P\left(w; w^{(t)}\right) = f\left(w^{(t)}\right) + \langle w - w^{(t)}, \nabla f\left(w^{(t)}\right)\rangle.$$

while hoping that this mechanism would push the produced chain $\left\{w^{(t)}\right\}$ towards the minimum because of (2.3). Hence we could recursively minimize the linear approximation (2.4) and the distance between the current state $w^{(t)}$ and the next w value to produce $w^{(t+1)}$; namely

(2.5)
$$\begin{aligned}
w^{(t+1)} &= \arg\min_{\forall w} \left(\frac{1}{2}\left\|w - w^{(t)}\right\|^2 + \eta P\left(w; w^{(t)}\right)\right) \\
&= \arg\min_{\forall w} \left(\frac{1}{2}\left\|w - w^{(t)}\right\|^2 + \eta\left(f\left(w^{(t)}\right) + \langle w - w^{(t)}, \nabla f\left(w^{(t+1)}\right)\rangle\right)\right) \\
&= w^{(t)} - \eta\nabla f\left(w^{(t)}\right)
\end{aligned}$$

where parameter $\eta > 0$ controls the trade off in (2.5).

Remark 7. Given T GD algorithm iterations, the output of GD can be (but not a exclusively),

(1) the average (after discarding the first few iterations of $w^{(t)}$ for stability reasons)

(2.6)
$$w_{\text{GD}}^{(T)} = \frac{1}{T} \sum_{t=1}^{T} w^{(t)}$$

(2) or the best value discovered
$$w_{\text{GD}}^{(T)} = \arg\min_{\forall w_t} \left(f\left(w^{(t)}\right) \right)$$

(3) or the last value discovered
$$w_{\text{GD}}^{(T)} = w^{(T)}$$

Note 8. GD output converges to a local minimum, $w_{\text{GD}}^{(T)} \to w_*$ (in some sense), under different sets of regularity conditions (some are weaker other stronger). Section 4 has a brief analysis.

Remark 9. The parameter η_t is called learning rate, step size, or gain. It determines the size of the steps GD takes to reach a (local) minimum. $\{\eta_t\}$ is a non-negative sequence and it is chosen by the practitioner. In principle, regularity conditions (Note 8) often imply restrictions on the decay of $\{\eta_t\}$ which guide the practitioner to parametrize it properly. Some popular choices of learning rate η_t are:

(1) <u>constant</u>; $\eta_t = \eta$, for where $\eta > 0$ is a small value. The rationale is that GD chain $\{w_t\}$ performs constant small steps towards the (local) minimum w_* and then oscillate around it.
(2) <u>decreasing and converging to zero</u>; $\eta_t \searrow$ with $\lim_{t\to\infty} \eta_t = 0$. E.g. $\eta_t = \left(\frac{C}{t}\right)^\varsigma$ where $\varsigma \in [0.5, 1]$ and $C > 0$. The rationale is that GD algorithm starts by performing larger steps (controlled by C) at the begining to explore the area for discovering possible minima. Also it reduces the size of those steps with the iterations (controled by ς) such that eventually when the chain $\{w_t\}$ is close to a possible minimum w_* value to converge and do not overshoot.
(3) <u>decreasing and converging to a tiny value</u> τ_*; $\eta_t \searrow$ with $\lim_{t\to\infty} \eta_t = \tau_*$ E.g. $\eta_t = \left(\frac{C}{t}\right)^\varsigma + \tau_*$ with $\varsigma \in (0.5, 1]$, $C > 0$, and $\tau_* \approx 0$. Same as previously, but the algorithm aims at oscillating around the detected local minimum.
(4) <u>constant until an iteration T_0 and then decreasing</u>; Eg $\eta_t = \left(\frac{C}{\max(t, T_0)}\right)^\varsigma$ with $\varsigma \in [0.5, 1]$ and $C > 0$, and $T_0 < T$. The rationale is that at the first stage of the iterations (when $t \le T_0$) the algorithm may need a constant large steps for a significant number of iterations T_0 in order to explore the domain; and hence in order for the chain $\{w_t\}$ to reach the area around the (local) minimum w_*. In the second stage, hoping that the chain $\{w_t\}$ may be in close proximity to the (local) minimum w_* the algorithm progressively performs smaller steps to converge towards the minimum w_*. The first stage ($t \le T_0$) is called burn-in; the values $\{w_t\}$ produced during the burn-in ($t \le T_0$) are are often discarded/ignored from the output of the GD algorithm.
   • Parameters $C, \varsigma, \tau_*, T_0$ may be chosen based on pilot runs.

Remark 10. There are several practical termination criteria that can be used in GD Algorithm 1(step 2). They aim to terminate the recursion in practice. Some popular termination criteria are

Created on 2023/02/26 at 16:53:13 by Georgios Karagiannis

(1) terminate when the gradient is sufficiently close to zero; i.e. if $\left\|\nabla f\left(w^{(t)}\right)\right\| \leq \epsilon$ for some pre-specified tiny $\epsilon > 0$ then STOP

(2) terminate when the chain $w^{(t)}$ does not change; i.e. if $\left\|w^{(t+1)} - w^{(t)}\right\| \leq \epsilon \left\|w^{(t)}\right\|$ for some pre-specified tiny $\epsilon > 0$ then STOP

(3) terminate when a pre-specified number of iterations T is performed; i.e. if $t \geq T$ then STOP

Here (1) may be deceive if the chain is in a flat area, (2) may be deceived if the learning rate become too small, (3) is obviously a last resort.

3. GD FOR NON-DIFFERENTIABLE FUNCTIONS (USING SUB-GRADIENTS)

Note 11. In several learning problems the function to be minimized is not differentiable. GD can be extended to address such problems with the use of subgradients.
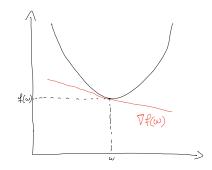
Definition 12. Vector v is called subgradient of a function $f : S \to \mathbb{R}$ at $w \in S$ if

$$(3.1) \qquad \forall u \in S, \ \ f(u) \geq f(w) + \langle u - w, v \rangle$$

Note 13. Essentially there may be more than one subgradients of the function at a specific point. As seen by (3.1), subgradients are the slopes of all the lines passing through the point $(w, f(w))$ and been under the function $f(\cdot)$.

(A) subgradients satisfying (3.1) in the non-differentiable case

(B) gradient satisfying the equality in (3.1) in the differentiable case

Notation 14. The set of subgradients of function $f : S \to \mathbb{R}$ at $w \in S$ is denoted by $\partial f(w)$.

Definition 15. The Gradient Descent algorithm using subgradients in non-differentiable cases, results by replacing the gradient $\nabla f\left(w^{(t)}\right)$ in (2.2) with any of subgradient v_t from the set of subgradients $\partial f\left(w^{(t)}\right)$ at $w^{(t)}$; namely

$$(3.2) \qquad w^{(t+1)} = w^{(t)} - \eta_t v_t; \ \text{ where } \ v_t \in \partial f\left(w^{(t)}\right)$$

3.1. **Construction of subgradient.**

Note 16. We discuss how to construct subgradients in practice.

Fact 17. *Some properties of subgradient sets that help for their construction*

(1) *If function $f : S \to \mathbb{R}$ is differentiable at w then the only subgradient of f at w is the gradient $\nabla f(w)$, and (3.1) is equality; i.e. $\partial f(w) = \{\nabla f(w)\}$.*

(2) *for constants α, β and convex function $f(\cdot)$, it is*

$$\partial (\alpha f(w) + \beta) = \alpha (\partial f(w)) = \{\alpha v : v \in \partial f(w),\}$$

(3) *for convex functions $f(\cdot)$ and $g(\cdot)$, it is*

$$\partial (f(w) + g(w)) = \partial f(w) + \partial g(w) = \{v + u : v \in \partial f(w),\, and\, u \in \partial g(w)\}$$

Example 18. Consider the function $f : \mathbb{R} \to \mathbb{R}_+$ with $f(w) = |w| = \begin{cases} w & w \geq 0 \\ -w & w < 0 \end{cases}$. Find the set

of subgradients $\partial f(w)$ for each $w \in \mathbb{R}$.

Solution. Using Fact 17, it is $\partial f(w) = 1$ for $w > 0$ and $\partial f(w) = -1$ for $w < 0$ as f is differentiable for $x \neq 0$. At $x = 0$, f is not differentiable; hence from condition (3.1) it is

$$\forall u \in \mathbb{R}, \;\; |u| \geq |0| + (u - 0)v$$

which is satisfied for $v \in [-1, 1]$. Hence,

$$\partial f(w) = \begin{cases} \{-1\} & , w < 0 \\ [-1, 1] & , w = 0 \\ \{1\} & , w > 0 \end{cases}$$

4. ANALYSIS OF GRADIENT DESCENT

Assumption 19. *For the sake of the analysis of the GD, let us consider:*

(1) *constant learning rate $\eta_t = \eta$,*
(2) *GD output $w_{GD}^{(T)} = \frac{1}{T} \sum_{t=1}^{T} w^{(t)}$*

Also as we will see, function $f(\cdot)$ should be convex and Lipschitz fin order for the following results to hold

Notation 20. w^* is the minimizer in (2.1).

Note 21. We consider Lemma 22 as given Fact.

No need to

Lemma 22. *Let $\{v_t; t = 1, ..., T\}$ be a sequence of vectors. Any algorithm with $w^{(1)} = 0$ and* memorize
$w^{(t+1)} = w^{(t)} - \eta v_t$ for $t = 1, ..., T$ satisfies

$$(4.1) \qquad \sum_{t=1}^{T} \langle w^{(t)} - w^*, v_t \rangle \leq \frac{\|w^*\|^2}{2\eta} + \frac{\eta}{2} \sum_{t=1}^{T} \|v_t\|^2$$

Proof. Omitted; see the 2nd reference. □

Note 23. To find an upper bound of the GD error, we try to bound the error $f\left(w_{\text{GD}}^{(T)}\right) - f\left(w^*\right)$ with purpose to use Lemma 22.

Proposition 24. *Consider the minimization problem (2.1). Given Assumptions 19, the error can be bounded as*

$$(4.2) \qquad f\left(w_{GD}^{(T)}\right) - f\left(w^*\right) \leq \frac{1}{T} \sum_{t=1}^{T} \langle w^{(t)} - w^*, v_t \rangle$$

where $v_t \in \partial f\left(w^{(t)}\right)$. If $f\left(\cdot\right)$ is differentiable then $v_t = f\left(w^{(t)}\right)$

Proof. It is[1]

$$f\left(w_{\text{GD}}^{(T)}\right) - f\left(w^*\right) = f\left(\frac{1}{T} \sum_{t=1}^{T} w_t\right) - f\left(w^*\right)$$

$$(4.3) \qquad\qquad\qquad \leq \frac{1}{T} \sum_{t=1}^{T} \left(f\left(w_t\right) - f\left(w^*\right)\right) \qquad \text{(by Jensen's inequality)}$$

$$\qquad\qquad\qquad \leq \frac{1}{T} \sum_{t=1}^{T} \langle w^{(t)} - w^*, \nabla f\left(w^{(t)}\right) \rangle \qquad \text{(by convexity of } f\left(\cdot\right)\text{)}$$

□

Note 25. The following provides conditions under which the gradient (or sub-gradient) is bounded. This is necessary in order to bound (4.2) with (4.1) in a meaningful manner.

Proposition 26. [2]*$f : S \to \mathbb{R}$ is ρ-Lipschitz over an open convex set S if and only if for all $w \in S$ and $v \in \partial f\left(w\right)$ it is $\|v\| \leq \rho$.*

Proof. $\implies$ Let $f : S \to \mathbb{R}$ be ρ-Lipschitz over convex set S, $w \in S$ and $v \in \partial f\left(w\right)$.

- Since S is open we get that there exist $\epsilon > 0$ such as $u := w + \epsilon \frac{v}{\|v\|}$ where $u \in S$. So $\langle u - w, v \rangle = \epsilon \|v\|$ and $\|u - w\| = \epsilon$.
- From the subgradient definition we get

$$f\left(u\right) - f\left(w\right) \geq \langle u - w, v \rangle = \epsilon \|v\|$$

- From the Lipschitzness of $f\left(\cdot\right)$ we get

$$f\left(u\right) - f\left(w\right) \leq \rho \|u - w\| = \rho\epsilon$$

Therefore $\|v\| \leq \rho$.

For $\impliedby$ see Exercise 4 in the Exercise sheet. □

---

[1]Jensen's inequality for convex $f\left(.\right)$ is $\text{E}\left(f\left(x\right)\right) \leq f\left(\text{E}\left(x\right)\right)$

[2]If this was a Homework there would be a Hint:

- If S is open there exist $\epsilon > 0$ such as $u = w + \epsilon \frac{v}{\|v\|}$ such as $u \in S$

Note 27. The following summaries Lemma 22, and Propositions 24, 26 with respect to the GD algorithm.

Proposition 28. *Let $f(\cdot)$ be a convex and ρ-Lipschitz function. If we run GD algorithm of f with learning rate $\eta > 0$ for T steps the output $w_{\mathrm{GD}}^{(T)} = \frac{1}{T}\sum_{t=1}^{T} w^{(t)}$ satisfies*

$$f\left(w_{\mathrm{GD}}^{(T)}\right) - f(w^*) \le \frac{\|w^*\|^2}{2\eta T} + \frac{\eta}{2}\frac{1}{T}\sum_{t=1}^{T}\left\|\nabla f\left(w^{(t)}\right)\right\|^2$$

where $\|\nabla f(\cdot)\| \le \rho$.

Solution. Straightforward from Lemma 22, and Propositions 24, 26.

Note 29. The following shows that a given learning rate depending on the iteration t, we can reduce the upper bound of the error as well as find the number of required iterations to achieve convergence.

Proposition 30. *(Cont Prop. 28)Let $f(\cdot)$ be a convex and ρ-Lipschitz function, and let $\mathcal{H} = \{w \in \mathbb{R} : \|w\| \le B\}$. Assume we run GD algorithm of $f(\cdot)$ with learning rate $\eta_t = \sqrt{\frac{B^2}{\rho^2 T}}$ for T steps, and output $w_{\mathrm{GD}}^{(T)} = \frac{1}{T}\sum_{t=1}^{T} w^{(t)}$. Then*

(1) *upper bound on the sub-optimality is*

(4.4) $$f\left(w_{\mathrm{GD}}^{(T)}\right) - f(w^*) \le \frac{B\rho}{\sqrt{T}}$$

(2) *a given level off accuracy ε such that $f\left(w_{GD}^{(T)}\right) - f(w^*) \le \varepsilon$ can be achieved after T iterations*

$$T \ge \frac{B^2\rho^2}{\varepsilon^2}.$$

Proof. Part 1 is a simple substitution from Proposition 28, and part 2 is implied from part 1. ☐

Note 31. The result on Proposition 30 heavily relies on setting suitable values for B and ρ which is rather a difficult task to be done in very complicated learning problems (e.g., learning a neural network).

Remark 32. The above results from the analysis of the GD also hold for the GD with subgradients; just replace $\nabla f(\cdot)$ with any v_t such that $v_t \in \partial f(\cdot)$.

5. Examples[3]

Example 33. Consider the simple Normal linear regression problem where the dataset $\{z_i = (y_i, x_i)\}_{i=1}^{n} \in \mathcal{D}$ is generated from a Normal data generating model

(5.1) $$\begin{pmatrix} y_i \\ x_i \end{pmatrix} \overset{\text{iid}}{\sim} \mathrm{N}\left(\begin{pmatrix} \mu_y \\ \mu_x \end{pmatrix}, \begin{bmatrix} \sigma_y^2 & \rho\sqrt{\sigma_y^2\sigma_x^2} \\ \rho\sqrt{\sigma_y^2\sigma_x^2} & \sigma_x^2 \end{bmatrix}\right)$$

---

[3]Code is available in `https://github.com/georgios-stats/Machine_Learning_and_Neural_Networks_III_Epiphany_2023/tree/main/Lecture_handouts/code/02.Gradient_descent/example_1.R`

for $i = 1, ..., n$. Consider a hypothesis space $\mathcal{H}$ of linear functions $h : \mathbb{R}^2 \to \mathbb{R}$ with $h(w) = w_1 + w_2 x$. The exact solution (which we pretend we do not know) is given as

$$\begin{pmatrix} w_1^* \\ w_2^* \end{pmatrix} = \begin{pmatrix} \mu_y - \rho \frac{\sigma_y}{\sigma_x} \mu_x \\ \rho \frac{\sigma_y}{\sigma_x} \end{pmatrix}.$$

To learn the optimal $w^* = (w_1^*, w_2^*)^\top$, we consider a loss $\ell \left(w, z_i = (x_i, y_i)^\top \right) = (y_i - [w_1 + w_2 x_i])^2$, which leads to the minimization problem

$$w^* = \arg \min_{\forall w} \left(\hat{R}_\mathcal{D}(w) \right) = \arg \min_{\forall w} \left(\frac{1}{n} \sum_{i=1}^{n} (y_i - w_1 - w_2 x_i)^2 \right)$$

The GD Algorithm 1 with learning rate η is

For $t = 1, 2, 3, ...$ iterate:

(1) compute

$$w^{(t+1)} = w^{(t)} - \eta v_t,$$

where $v_t = \begin{pmatrix} 2w_1^{(t)} + 2w_2^{(t)} \bar{x} - 2\bar{y} \\ 2w_1^{(t)} \bar{x} + 2w_2^{(t)} \overline{x^2} - 2y^\top x \end{pmatrix}$

(2) terminate if a termination criterion is satisfied, e.g.

If $t \geq T$ then STOP

This is because $\hat{R}_\mathcal{D}(w)$ is differentiable in $\mathbb{R}^2$ so $\partial \hat{R}_\mathcal{D}(w) = \left\{ \nabla \hat{R}_\mathcal{D}(w) \right\}$ and because

$$\nabla \hat{R}_\mathcal{D}(w) = \begin{pmatrix} \frac{\mathrm{d}}{\mathrm{d}w_1} \hat{R}_\mathcal{D}(w) \\ \frac{\mathrm{d}}{\mathrm{d}w_2} \hat{R}_\mathcal{D}(w) \end{pmatrix} = ... = \begin{pmatrix} 2w_1^{(t)} + 2w_2^{(t)} \bar{x} - 2\bar{y} \\ 2w_1^{(t)} \bar{x} + 2w_2^{(t)} \overline{x^2} - 2y^\top x \end{pmatrix}$$

Consider data size $n = 100$, and parameters $\rho = 0.2$, $\sigma_y^2 = 1$ and $\sigma_x^2 = 1$. Then the real value (5.2) that I need to learn equals to $w^* = (0, 1)^\top$. Consider a GD seed $w_0 = (2, -2)$, and total number of iterations $T = 1000$.

Figures 5.1a, 5.1b, and 5.1c present trace plots of the chain $\left\{ (w^{(t)}) \right\}$ and error $\hat{R}_\mathcal{D}(w^{(t)}) - \hat{R}_\mathcal{D}(w^*)$ produced by running GD for $T = 1000$ total iterations and for different (each time) constant learning rates $\eta \in \{0.01, 0.02, 0.05, 0.99\}$. We observe that the larger learning rates under consideration were able to converge faster to the minimum w^*. This is because they perform larger steps and can learn faster -this is not a panacea.

Figures 5.1d, 5.1e, and 5.1f present trace plots of the chain $\left\{ (w^{(t)}) \right\}$, and of the error $\hat{R}_\mathcal{D}(w^{(t)}) - \hat{R}_\mathcal{D}(w^*)$ produced by running GD for $T = 1000$ total iterations and for learning rate $\eta = 1.0$ (previously considered) and a very big learning rate $\eta = 3.0$. We observe that the very big learning rate $\eta = 3.0$ presents slower convergence to the minimum w^*. This is because it creates unreasonably big steps in (2.2) that the produced chain overshoots the global minimum. See the cartoon in Figures 5.1a and 5.1b.

(A) $\left\{\left(w_1^{(t)}\right)\right\}$

(B) $\left\{\left(w_2^{(t)}\right)\right\}$

(C) $\hat{R}_{\mathcal{D}}\left(w^{(t)}\right) - \hat{R}_{\mathcal{D}}\left(w^*\right)$

(D) $\left\{\left(w_1^{(t)}\right)\right\}$

(E) $\left\{\left(w_2^{(t)}\right)\right\}$

(F) $\hat{R}_{\mathcal{D}}\left(w^{(t)}\right) - \hat{R}_{\mathcal{D}}\left(w^*\right)$

(A) Unnecessarily large learning rate

(B) Unnecessarily small learning rate

Created on 2023/02/26 at 16:53:13   by Georgios Karagiannis